

Fast low-rank metric learning


Dan-Theodor Oneață

July 25, 2011

k nearest neighbours

- ▶ Simple, yet powerful classifier.
- ▶ Euclidean distance:


$$d(\mathbf{x}_i, \mathbf{x}_j) = \sqrt{(\mathbf{x}_i - \mathbf{x}_j)^T (\mathbf{x}_i - \mathbf{x}_j)}$$



k nearest neighbours

- ▶ Simple, yet powerful classifier.
- ▶ Euclidean distance:


$$d(\mathbf{x}_i, \mathbf{x}_j) = \sqrt{(\mathbf{x}_i - \mathbf{x}_j)^T (\mathbf{x}_i - \mathbf{x}_j)}$$



k nearest neighbours

- ▶ Simple, yet powerful classifier.
- ▶ Euclidean distance:

$$d(\mathbf{x}_i, \mathbf{x}_j) = \sqrt{(\mathbf{x}_i - \mathbf{x}_j)^T (\mathbf{x}_i - \mathbf{x}_j)}$$

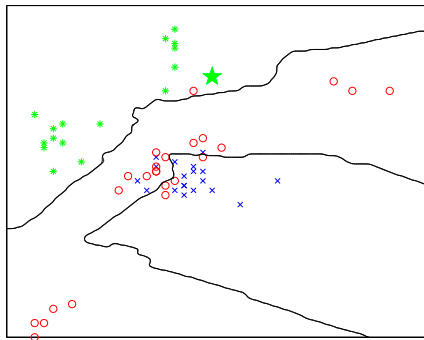


Decision boundaries for $k = 1$

k nearest neighbours

- ▶ Simple, yet powerful classifier.
- ▶ Euclidean distance:

$$d(\mathbf{x}_i, \mathbf{x}_j) = \sqrt{(\mathbf{x}_i - \mathbf{x}_j)^T (\mathbf{x}_i - \mathbf{x}_j)}$$



Decision boundaries for $k = 7$

k nearest neighbours

- ▶ Simple, yet powerful classifier.
- ▶ Euclidean distance:


$$d(\mathbf{x}_i, \mathbf{x}_j) = \sqrt{(\mathbf{x}_i - \mathbf{x}_j)^T (\mathbf{x}_i - \mathbf{x}_j)}$$



k nearest neighbours

- ▶ Simple, yet powerful classifier.
- ▶ Euclidean distance:

$$d(\mathbf{x}_i, \mathbf{x}_j) = \sqrt{(\mathbf{x}_i - \mathbf{x}_j)^T (\mathbf{x}_i - \mathbf{x}_j)}$$



Face recognition

k nearest neighbours

- ▶ Simple, yet powerful classifier.
- ▶ Euclidean distance:

$$d(\mathbf{x}_i, \mathbf{x}_j) = \sqrt{(\mathbf{x}_i - \mathbf{x}_j)^T (\mathbf{x}_i - \mathbf{x}_j)}$$



Expression recognition

Neighbourhood component analysis


(Goldberger et al, 2004)

- Learns a Mahalanobis metric

$$d_{\mathbf{S}}(\mathbf{x}_i, \mathbf{x}_j) = \sqrt{(\mathbf{x}_i - \mathbf{x}_j)^T \mathbf{S} (\mathbf{x}_i - \mathbf{x}_j)}$$

- Equivalent to a linear transformation:

$$d_{\mathbf{S}}(\mathbf{x}_i, \mathbf{x}_j) = d_{\mathbf{I}}(\mathbf{A}\mathbf{x}_i, \mathbf{A}\mathbf{x}_j)$$



Mahalanobis metric

Neighbourhood component analysis


(Goldberger et al, 2004)

- Learns a Mahalanobis metric

$$d_{\mathbf{S}}(\mathbf{x}_i, \mathbf{x}_j) = \sqrt{(\mathbf{x}_i - \mathbf{x}_j)^T \mathbf{S} (\mathbf{x}_i - \mathbf{x}_j)}$$

- Equivalent to a linear transformation:

$$d_{\mathbf{S}}(\mathbf{x}_i, \mathbf{x}_j) = d_{\mathbf{I}}(\mathbf{A}\mathbf{x}_i, \mathbf{A}\mathbf{x}_j)$$



Euclidean metric in $\mathbf{A}\mathbf{X}$

Neighbourhood component analysis


(Goldberger et al, 2004)

1. Find \mathbf{A} that maximizes leave-one-out cross-validation score.
2. Soft version:

$$p(\mathbf{x}_i \in \text{class } c) = \frac{\sum_{j \in c} \exp\{-d_{\mathbf{S}}^2(\mathbf{x}_i, \mathbf{x}_j)\}}{\sum_k \exp\{-d_{\mathbf{S}}^2(\mathbf{x}_i, \mathbf{x}_k)\}}$$

Maximize $f(\mathbf{A})$

$$= \sum_i p(\mathbf{x}_i \in \text{true class of } \mathbf{x}_i).$$



Neighbourhood component analysis


(Goldberger et al, 2004)

1. Find \mathbf{A} that maximizes leave-one-out cross-validation score.
2. Soft version:

$$p(\mathbf{x}_i \in \text{class } c) = \frac{\sum_{j \in c} \exp\{-d_{\mathbf{S}}^2(\mathbf{x}_i, \mathbf{x}_j)\}}{\sum_k \exp\{-d_{\mathbf{S}}^2(\mathbf{x}_i, \mathbf{x}_k)\}}$$


Maximize $f(\mathbf{A})$

$$= \sum_i p(\mathbf{x}_i \in \text{true class of } \mathbf{x}_i).$$




Optimizing $f(\mathbf{A})$

- Use $\nabla_{\mathbf{A}} f(\mathbf{A})$ for an optimization algorithm: *e.g.*, gradient ascent, conjugate gradients.
- How to initialise? Use random \mathbf{A} or most discriminative projections given by PCA, LDA or logistic regression.



Optimizing $f(\mathbf{A})$

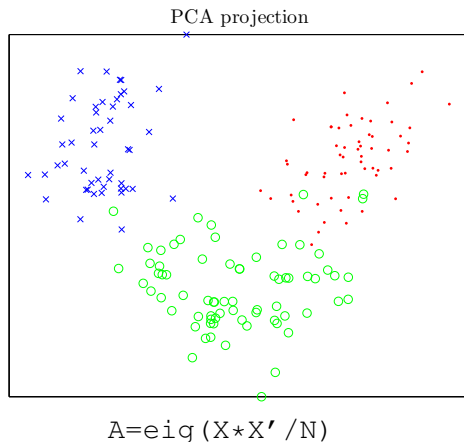
- Use $\nabla_{\mathbf{A}} f(\mathbf{A})$ for an optimization algorithm: *e.g.*, gradient ascent, conjugate gradients.
- How to initialise? Use random \mathbf{A} or most discriminative projections given by PCA, LDA or logistic regression.



`A=maximize('nca',A)`


Optimizing $f(\mathbf{A})$

- Use $\nabla_{\mathbf{A}} f(\mathbf{A})$ for an optimization algorithm: *e.g.*, gradient ascent, conjugate gradients.
- How to initialise? Use random \mathbf{A} or most discriminative projections given by PCA, LDA or logistic regression.



Optimizing $f(\mathbf{A})$

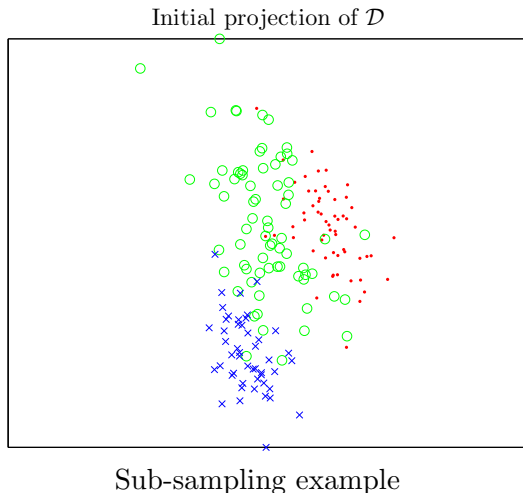
- Use $\nabla_{\mathbf{A}} f(\mathbf{A})$ for an optimization algorithm: *e.g.*, gradient ascent, conjugate gradients.
- How to initialise? Use random \mathbf{A} or most discriminative projections given by PCA, LDA or logistic regression.



$\mathbf{A} = \text{maximize}('nca', \mathbf{A})$


Speeding up the computations

1. Sub-sample the data set.
2. Use mini-batches:
 - ▶ Choose them randomly
 - ▶ Use cheap clustering method.




Speeding up the computations

1. Sub-sample the data set.
2. Use mini-batches:
 - ▶ Choose them randomly
 - ▶ Use cheap clustering method.




Speeding up the computations

1. Sub-sample the data set.
2. Use mini-batches:
 - ▶ Choose them randomly
 - ▶ Use cheap clustering method.




Speeding up the computations

1. Sub-sample the data set.
2. Use mini-batches:
 - ▶ Choose them randomly
 - ▶ Use cheap clustering method.



Speeding up the computations


1. Sub-sample the data set.
2. Use mini-batches:
 - ▶ Choose them randomly
 - ▶ Use cheap clustering method.



Recursive projection clustering

Speeding up the computations


1. Sub-sample the data set.
2. Use mini-batches:
 - ▶ Choose them randomly
 - ▶ Use cheap clustering method.



Recursive projection clustering

Speeding up the computations


1. Sub-sample the data set.
2. Use mini-batches:
 - ▶ Choose them randomly
 - ▶ Use cheap clustering method.



Recursive projection clustering

Approximate computations


- Each contribution $p_{ij} \propto \exp\{-d^2(\mathbf{A}\mathbf{x}_i, \mathbf{A}\mathbf{x}_j)\}$.
- Cast NCA into class conditional kernel density estimation problem: $p(\mathbf{x}_i|c)$.
- Use k -d trees for fast density estimation.



How p_{ij} varies during training

Approximate computations


- ▶ Each contribution $p_{ij} \propto \exp\{-d^2(\mathbf{A}\mathbf{x}_i, \mathbf{A}\mathbf{x}_j)\}$.
- ▶ Cast NCA into class conditional kernel density estimation problem: $p(\mathbf{x}_i|c)$.
- ▶ Use k -d trees for fast density estimation.



k -d tree example

Approximate computations

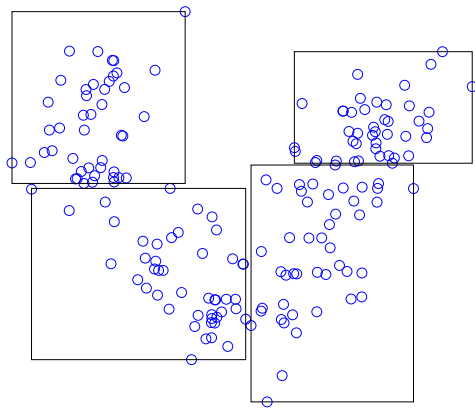
- Each contribution $p_{ij} \propto \exp\{-d^2(\mathbf{A}\mathbf{x}_i, \mathbf{A}\mathbf{x}_j)\}$.
- Cast NCA into class conditional kernel density estimation problem: $p(\mathbf{x}_i|c)$.
- Use k -d trees for fast density estimation.



k -d tree example

Approximate computations


- ▶ Each contribution $p_{ij} \propto \exp\{-d^2(\mathbf{A}\mathbf{x}_i, \mathbf{A}\mathbf{x}_j)\}$.
- ▶ Cast NCA into class conditional kernel density estimation problem: $p(\mathbf{x}_i|c)$.
- ▶ Use k -d trees for fast density estimation.



k -d tree example

Approximate computations


- ▶ Each contribution $p_{ij} \propto \exp\{-d^2(\mathbf{A}\mathbf{x}_i, \mathbf{A}\mathbf{x}_j)\}$.
- ▶ Cast NCA into class conditional kernel density estimation problem: $p(\mathbf{x}_i|c)$.
- ▶ Use k -d trees for fast density estimation.



k -d tree example

Approximate computations


- ▶ Each contribution $p_{ij} \propto \exp\{-d^2(\mathbf{A}\mathbf{x}_i, \mathbf{A}\mathbf{x}_j)\}$.
- ▶ Cast NCA into class conditional kernel density estimation problem: $p(\mathbf{x}_i|c)$.
- ▶ Use k -d trees for fast density estimation.



k -d tree example

Approximate computations


- ▶ Each contribution $p_{ij} \propto \exp\{-d^2(\mathbf{A}\mathbf{x}_i, \mathbf{A}\mathbf{x}_j)\}$.
- ▶ Cast NCA into class conditional kernel density estimation problem: $p(\mathbf{x}_i|c)$.
- ▶ Use k -d trees for fast density estimation.



CC-KDE example: $p(\mathbf{x}_i|c = \text{red})$

Approximate computations


- Each contribution $p_{ij} \propto \exp\{-d^2(\mathbf{A}\mathbf{x}_i, \mathbf{A}\mathbf{x}_j)\}$.
- Cast NCA into class conditional kernel density estimation problem: $p(\mathbf{x}_i|c)$.
- Use k -d trees for fast density estimation.



CC-KDE example: $p(\mathbf{x}_i|c = \text{red})$

Approximate computations


- Each contribution $p_{ij} \propto \exp\{-d^2(\mathbf{A}\mathbf{x}_i, \mathbf{A}\mathbf{x}_j)\}$.
- Cast NCA into class conditional kernel density estimation problem: $p(\mathbf{x}_i|c)$.
- Use k -d trees for fast density estimation.



CC-KDE example: $p(\mathbf{x}_i|c = \text{red})$

Approximate computations


- ▶ Each contribution $p_{ij} \propto \exp\{-d^2(\mathbf{A}\mathbf{x}_i, \mathbf{A}\mathbf{x}_j)\}$.
- ▶ Cast NCA into class conditional kernel density estimation problem: $p(\mathbf{x}_i|c)$.
- ▶ Use k -d trees for fast density estimation.



CC-KDE example: $p(\mathbf{x}_i|c = \text{red})$

Approximate computations


- ▶ Each contribution $p_{ij} \propto \exp\{-d^2(\mathbf{A}\mathbf{x}_i, \mathbf{A}\mathbf{x}_j)\}$.
- ▶ Cast NCA into class conditional kernel density estimation problem: $p(\mathbf{x}_i|c)$.
- ▶ Use k -d trees for fast density estimation.



CC-KDE example: $p(\mathbf{x}_i|c = \text{red})$

Approximate computations


- ▶ Each contribution $p_{ij} \propto \exp\{-d^2(\mathbf{A}\mathbf{x}_i, \mathbf{A}\mathbf{x}_j)\}$.
- ▶ Cast NCA into class conditional kernel density estimation problem: $p(\mathbf{x}_i|c)$.
- ▶ Use k -d trees for fast density estimation.



CC-KDE example: $p(\mathbf{x}_i|c = \text{red})$

Approximate computations


- Each contribution $p_{ij} \propto \exp\{-d^2(\mathbf{A}\mathbf{x}_i, \mathbf{A}\mathbf{x}_j)\}$.
- Cast NCA into class conditional kernel density estimation problem: $p(\mathbf{x}_i|c)$.
- Use k -d trees for fast density estimation.



CC-KDE example: $p(\mathbf{x}_i|c = \text{red})$

Approximate computations


- ▶ Each contribution $p_{ij} \propto \exp\{-d^2(\mathbf{A}\mathbf{x}_i, \mathbf{A}\mathbf{x}_j)\}$.
- ▶ Cast NCA into class conditional kernel density estimation problem: $p(\mathbf{x}_i|c)$.
- ▶ Use k -d trees for fast density estimation.



CC-KDE example: $p(\mathbf{x}_i|c = \text{red})$

Approximate computations


- Each contribution $p_{ij} \propto \exp\{-d^2(\mathbf{A}\mathbf{x}_i, \mathbf{A}\mathbf{x}_j)\}$.
- Cast NCA into class conditional kernel density estimation problem: $p(\mathbf{x}_i|c)$.
- Use k -d trees for fast density estimation.



CC-KDE example: $p(\mathbf{x}_i|c = \text{red})$

Approximate computations


- ▶ Each contribution $p_{ij} \propto \exp\{-d^2(\mathbf{A}\mathbf{x}_i, \mathbf{A}\mathbf{x}_j)\}$.
- ▶ Cast NCA into class conditional kernel density estimation problem: $p(\mathbf{x}_i|c)$.
- ▶ Use k -d trees for fast density estimation.



CC-KDE example: $p(\mathbf{x}_i|c = \text{red})$

Approximate computations

- Each contribution $p_{ij} \propto \exp\{-d^2(\mathbf{A}\mathbf{x}_i, \mathbf{A}\mathbf{x}_j)\}$.
- Cast NCA into class conditional kernel density estimation problem: $p(\mathbf{x}_i|c)$.
- Use k -d trees for fast density estimation.



CC-KDE example: $p(\mathbf{x}_i|c = \text{red})$

Work done

- ▶ Introduction
- ▶ Background
 - ▶ Theoretical background
 - ▶ Related methods
- ▶ Neighbourhood component analysis
 - ▶ General presentation
 - ▶ Practical issues
 - ▶ Interpreting NCA as class-conditional density estimation problem
- ▶ Reducing computational cost
 - ▶ Sub-sampling
 - ▶ Mini-batches
 - ▶ Stochastic learning
 - ▶ Approximate computations
 - ▶ Exact computations using compact support kernels
 - ▶ NCA with compact support kernels and background distribution
- ▶ Evaluation (on large data sets)
- ▶ Conclusions

Remaining work

- ▶ Introduction
- ▶ Background
 - ▶ Theoretical background
 - ▶ Related methods
- ▶ Neighbourhood component analysis
 - ▶ General presentation
 - ▶ Practical issues
 - ▶ Interpreting NCA as class-conditional density estimation problem
- ▶ Reducing computational cost
 - ▶ Sub-sampling
 - ▶ Mini-batches
 - ▶ Stochastic learning
 - ▶ Approximate computations
 - ▶ Exact computations using compact support kernels
 - ▶ NCA with compact support kernels and background distribution
- ▶ Evaluation (on large data sets)
- ▶ Conclusions

Remaining work

- ▶ Introduction
- ▶ Background
 - ▶ Theoretical background
 - ▶ Related methods
- ▶ Neighbourhood component analysis
 - ▶ General presentation
 - ▶ Practical issues
 - ▶ Interpreting NCA as class-conditional density estimation problem
- ▶ Reducing computational cost
 - ▶ Sub-sampling
 - ▶ Mini-batches
 - ▶ Stochastic learning
 - ▶ Approximate computations
 - ▶ Exact computations using compact support kernels
 - ▶ NCA with compact support kernels and background distribution
- ▶ Evaluation (on large data sets)
- ▶ Conclusions

THANK YOU!