

Fast low-rank metric learning


Dan-Theodor Oneață

July 25, 2011

k nearest neighbours

- ▶ Simple, yet powerful classifier.
- ▶ Euclidean distance:


$$d(\mathbf{x}_i, \mathbf{x}_j) = \sqrt{(\mathbf{x}_i - \mathbf{x}_j)^T (\mathbf{x}_i - \mathbf{x}_j)}$$



k nearest neighbours

- ▶ Simple, yet powerful classifier.
- ▶ Euclidean distance:


$$d(\mathbf{x}_i, \mathbf{x}_j) = \sqrt{(\mathbf{x}_i - \mathbf{x}_j)^T (\mathbf{x}_i - \mathbf{x}_j)}$$



k nearest neighbours

- ▶ Simple, yet powerful classifier.
- ▶ Euclidean distance:

$$d(\mathbf{x}_i, \mathbf{x}_j) = \sqrt{(\mathbf{x}_i - \mathbf{x}_j)^T (\mathbf{x}_i - \mathbf{x}_j)}$$




$k = 1$

k nearest neighbours

- ▶ Simple, yet powerful classifier.
- ▶ Euclidean distance:

$$d(\mathbf{x}_i, \mathbf{x}_j) = \sqrt{(\mathbf{x}_i - \mathbf{x}_j)^T (\mathbf{x}_i - \mathbf{x}_j)}$$



$k = 7$

k nearest neighbours

- ▶ Simple, yet powerful classifier.
- ▶ Euclidean distance:


$$d(\mathbf{x}_i, \mathbf{x}_j) = \sqrt{(\mathbf{x}_i - \mathbf{x}_j)^T (\mathbf{x}_i - \mathbf{x}_j)}$$



k nearest neighbours

- ▶ Simple, yet powerful classifier.
- ▶ Euclidean distance:

$$d(\mathbf{x}_i, \mathbf{x}_j) = \sqrt{(\mathbf{x}_i - \mathbf{x}_j)^T (\mathbf{x}_i - \mathbf{x}_j)}$$



Face recognition

k nearest neighbours

- ▶ Simple, yet powerful classifier.
- ▶ Euclidean distance:

$$d(\mathbf{x}_i, \mathbf{x}_j) = \sqrt{(\mathbf{x}_i - \mathbf{x}_j)^T (\mathbf{x}_i - \mathbf{x}_j)}$$



Expression recognition


Neighbourhood component analysis

- Learns a Mahalanobis metric

$$d_{\mathbf{S}}(\mathbf{x}_i, \mathbf{x}_j) = \sqrt{(\mathbf{x}_i - \mathbf{x}_j)^T \mathbf{S} (\mathbf{x}_i - \mathbf{x}_j)}$$

- Equivalent to a linear transformation:

$$d_{\mathbf{S}}(\mathbf{x}_i, \mathbf{x}_j) = d_{\mathbf{I}}(\mathbf{A}\mathbf{x}_i, \mathbf{A}\mathbf{x}_j)$$




Neighbourhood component analysis

- Learns a Mahalanobis metric

$$d_{\mathbf{S}}(\mathbf{x}_i, \mathbf{x}_j) = \sqrt{(\mathbf{x}_i - \mathbf{x}_j)^T \mathbf{S} (\mathbf{x}_i - \mathbf{x}_j)}$$

- Equivalent to a linear transformation:

$$d_{\mathbf{S}}(\mathbf{x}_i, \mathbf{x}_j) = d_{\mathbf{I}}(\mathbf{A}\mathbf{x}_i, \mathbf{A}\mathbf{x}_j)$$



Neighbourhood component analysis

1. Find \mathbf{S} that maximizes leave-one-out cross-validation score.


2. Soft version:

$$p(\mathbf{x}_i \in \text{class } c) = \frac{\sum_{j \in c} \exp\{-d_{\mathbf{S}}(\mathbf{x}_i, \mathbf{x}_j)\}}{\sum_k \exp\{-d_{\mathbf{S}}(\mathbf{x}_i, \mathbf{x}_k)\}}$$

$$\text{Maximize } f(\mathbf{S}) = \sum_i p(\mathbf{x}_i \in \text{true class of } \mathbf{x}_i).$$

Optimizing $f(\mathbf{S})$

- Use $\nabla_{\mathbf{S}} f(\mathbf{S})$ for an optimization algorithm: *e.g.*, gradient ascent, conjugate gradients.
- How to initialise? Use random \mathbf{S} or most discriminative projections given by PCA, LDA or logistic regression.




$\mathbf{A} = \text{randn}(d, D)$

Data set: wine.

Optimizing $f(\mathbf{S})$

- ▶ Use $\nabla_{\mathbf{S}} f(\mathbf{S})$ for an optimization algorithm: *e.g.*, gradient ascent, conjugate gradients.
- ▶ How to initialise? Use random \mathbf{S} or most discriminative projections given by PCA, LDA or logistic regression.




`A=minimize('nca',A)`

Data set: wine.

Optimizing $f(\mathbf{S})$

- Use $\nabla_{\mathbf{S}} f(\mathbf{S})$ for an optimization algorithm: *e.g.*, gradient ascent, conjugate gradients.
- How to initialise? Use random \mathbf{S} or most discriminative projections given by PCA, LDA or logistic regression.




$$\mathbf{A} = \text{eig}(\mathbf{X} \star \mathbf{X}' / N)$$

Data set: wine.

Optimizing $f(\mathbf{S})$

- Use $\nabla_{\mathbf{S}} f(\mathbf{S})$ for an optimization algorithm: *e.g.*, gradient ascent, conjugate gradients.
- How to initialise? Use random \mathbf{S} or most discriminative projections given by PCA, LDA or logistic regression.



`A=minimize('nca',A)`

Data set: wine.

Speeding the computations

1. Sub-sample the data set.
2. Use mini-batches:
 - ▶ Choose them randomly
 - ▶ Use cheap clustering method.

Speeding the computations

1. Sub-sample the data set.
2. Use mini-batches:
 - ▶ Choose them randomly
 - ▶ Use cheap clustering method.

Approximate computations

Exact computations

Future work

Conclusions